

# Meteorological Super-Resolution vs Wind Representations

Gruppo 21  
Marzia De Maina, Matteo Galiazzo, Federica Santisi

*Alma Mater Studiorum - Università di Bologna*

*Dipartimento di Informatica - Scienza e Ingegneria (DISI)*

- Evaluate the impact of **different wind representations** on neural-based super-resolution for meteorological fields.
- Compare two **wind encodings**:
  - Orthogonal components  $u_{10}, v_{10}$
  - Polar form: speed and direction (degrees from north)
- Use **ERA5 (30 km)** and **VHR-REA (2.2 km)** datasets for training and testing.
- Assess which representation yields better performance and qualitative reconstruction.

## Cartesian Components

- $u_{10}, v_{10}$ : zonal and meridional wind at 10m height
- Directly match ERA5 outputs

## Polar Encoding

- speed =  $\sqrt{u_{10}^2 + v_{10}^2}$
- direction =  $(180 + \frac{180}{\pi} \arctan 2(-u_{10}, -v_{10})) \text{ mod } 360$

**Motivation:** Representation may affect neural network learning and reconstruction quality.

## Super-Resolution

Upscaling low-resolution fields to high-resolution targets using deep learning. This technique involves reconstructing high-resolution images from coarser data, with the aim of capturing finer and more complex atmospheric structures.

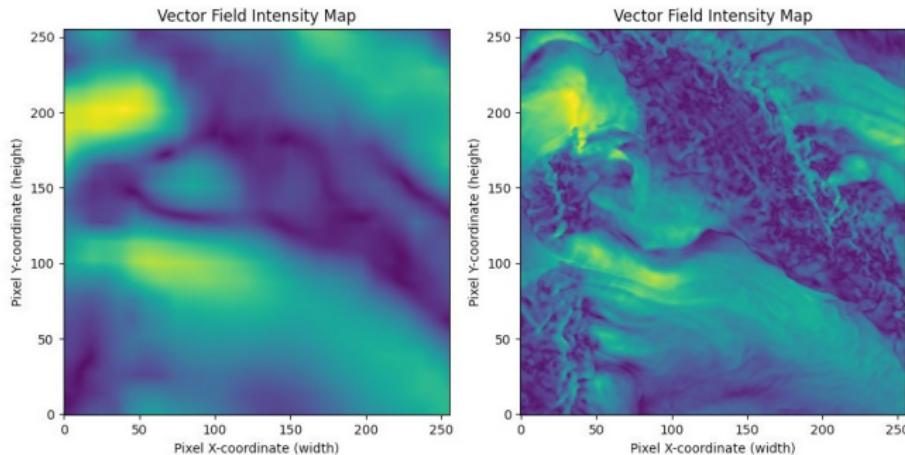


Figure: Super-resolution in the best case

## U-Net Architecture

Encoder-decoder with skip connections for image-to-image tasks.

- Residual blocks to ease training.
- Multiscale features preserved via concatenation.

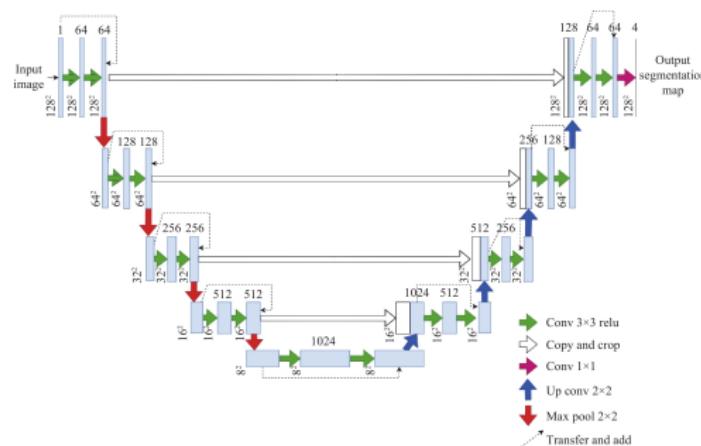


Figure: U-Net Standard Architecture

For the purpose of this project, the wind field representations are derived from the ERA5 and VHR–REA datasets.

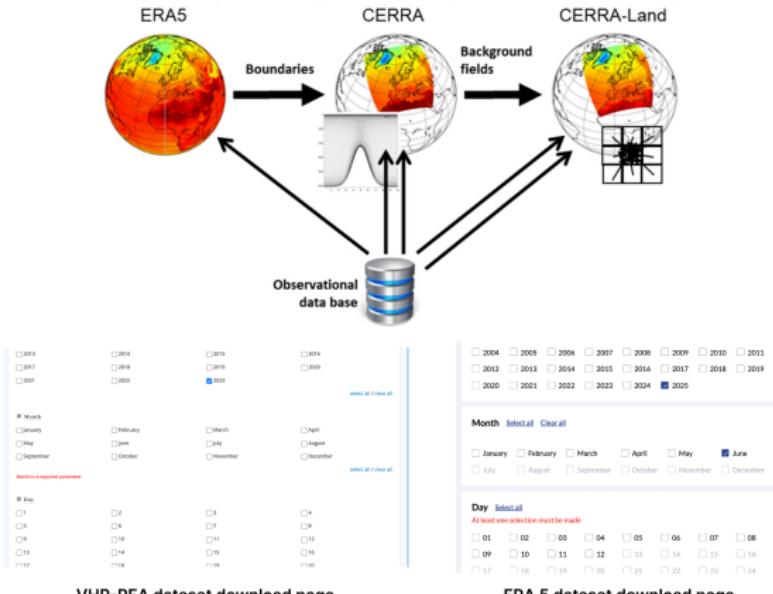
- **ERA5** is a global reanalysis dataset developed by ECMWF, providing hourly data since 1950 at a  $0.25^\circ$  ( $\sim 31$  km) resolution.
- **VHR–REA** is a downscaled regional reanalysis for Italy based on the COSMO model, with a resolution of 2.2 km.

Both datasets are aligned temporally (06, 12, 18, 00 UTC) and spatially through reprojection.

ERA5 acts as the low-resolution input; VHR–REA is used as the high-resolution target for training.

# From Global to Surface Reanalysis

Global Reanalysis → Regional Reanalysis → Surface Reanalysis



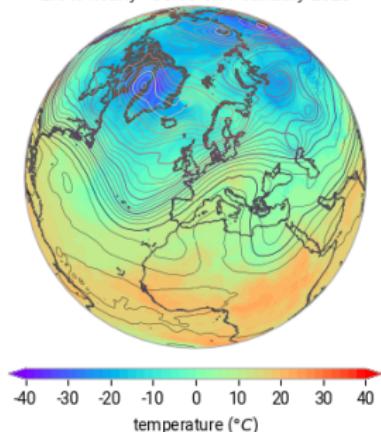
# Datasets - ERA5

**ERA5** is a global reanalysis dataset by ECMWF under the Copernicus Climate Change Service (C3S). It offers hourly data from 1950 at  $0.25^\circ$  resolution, replacing the older ERA-Interim.

Based on 4D-Var data assimilation, ERA5 integrates satellite, radiosonde, buoy, and aircraft observations, covering atmosphere, land, and ocean.

In this project, we use surface wind components  $u_{10}$  and  $v_{10}$  as low-resolution inputs for super-resolution modeling.

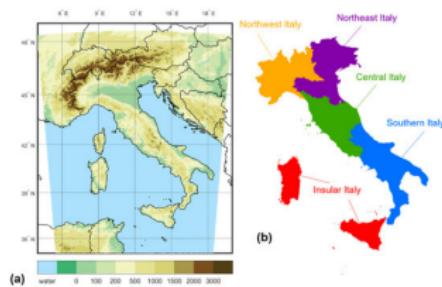
850 hPa temperature and 500 hPa geopotential  
ERA5 hourly - 00:00 on 1 January 2023



**VHR-REA (Very High Resolution ReAnalysis)** is produced by CMCC via dynamic downscaling of ERA5 using the COSMO model. It reaches a spatial resolution of 2.2 km, resolving convection in complex terrain.

This dataset is designed for high-resolution applications such as hydrology, renewable energy, and urban climate.

- Used as high-resolution target for training.
- Temporal coverage: 4 times daily for one full year.



**Figure:** VHR-REA spatial coverage over Northern Italy

# ERA5 & VHR–REA: Regridding Pipeline

## ERA5 → VHR Reprojection Strategy

- ERA5 loaded with Dask chunking
- Target grid built from concatenated VHR quarters
- Bilinear regridding using xesmf
- Reuse of interpolation weights (caching)
- Post-rechunking to control memory usage

## Data Output & Performance

- Lazy NetCDF writing with Dask
- Compute triggered with progress bar
- Output aligned with VHR spatial resolution

```
# Step 1: Load ERA5 with chunks
era5 = xr.open_dataset("era5.nc",
                      chunks={"time": 10, "lat": 300, "lon": 300})

# Step 2: Load and combine VHR target grid
vhr = xr.concat([
    xr.open_dataset("vhr_q1.nc"),
    xr.open_dataset("vhr_q2.nc")
], dim="time")

# Step 3: Create bilinear regridder
regridder = xe.Regridder(
    era5, vhr, method="bilinear",
    filename="weights.nc")

# Step 4: Apply regridding + rechunk
output = regridder(era5).chunk({
    "time": 10, "rlat": 100, "rlon": 100})

# Step 5: Save with progress bar
with ProgressBar():
    output.to_netcdf("regridded.nc")
```

# Preprocessing

## Spatial Alignment & Patch Extraction

- Center crop to  $224 \times 224$  pixel patches
- Temporal synchronization verification
- Channel-wise data stacking ( $u10, v10$  components)
- Custom PyTorch Dataset implementation

## Data Pipeline Architecture

- ItalyWeatherDataset class for paired data loading
- Automatic temporal alignment checking
- Flexible normalizer integration
- Memory-efficient batch processing

## Quality Assurance

- Index bounds verification
- Consistent tensor dimensions
- Error handling for misaligned datasets

```
def extract_region(self, data_tensor):
    patch_h, patch_w = (224, 224)
    _, h, w = data_tensor.shape
    x = (w - patch_w) // 2
    y = (h - patch_h) // 2
    return data_tensor[:, y:y+patch_h,
                      x:x+patch_w]

def __getitem__(self, idx):
    era_slice = self.era5_dataset.
               isel(valid_time=idx)
    era_arrays = [era_slice[v].
                  values
                  for v in self.
                  ERA5_VARIABLES]
    era_tensor = torch.from_numpy(
        np.stack(era_arrays, axis=0)
        .float())
    era_tensor = self.extract_region
    (era_tensor)

    if self.era5_normalizer:
        era_tensor = self.
                    era5_normalizer.
                    normalize(era_tensor)
```

# Normalization & Coordinate Transformation

## MinMax Normalization Strategy

- Training-based statistics computation
- Per-channel normalization across spatial dimensions
- Feature range scaling to [0, 1]
- Epsilon handling for division-by-zero cases
- Serializable normalizer objects

## Coordinate System Transformation

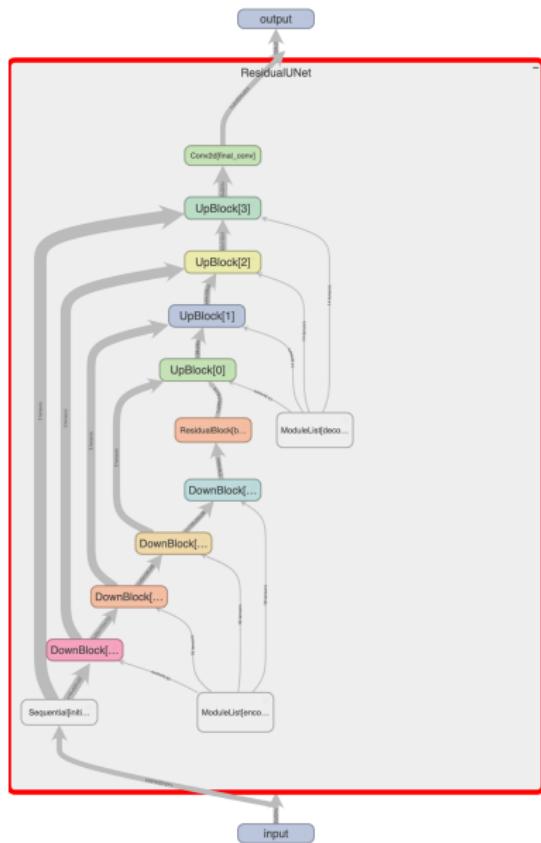
- Option 1: Cartesian coordinates ( $u$ ,  $v$ )
- Option 2: Polar coordinates (magnitude, direction)
  - $magnitude = \sqrt{u^2 + v^2}$
  - $direction = 180 + \frac{180}{\pi} \arctan 2(-u, -v) \bmod 360$

## Dataset Management

- 80-20 train-test split with fixed seed (42)
- Separate normalization preservation
- Batch processing with memory optimization

```
class MinMaxNormalizer:  
    def compute_stats(self, data_tensor):  
        # Min/max per channel across (N,H,W)  
        dims  
        self.min_val = data_tensor.amin(  
            dim=(0, 2, 3), keepdim=True)  
        self.max_val = data_tensor.amax(  
            dim=(0, 2, 3), keepdim=True)  
  
        # Handle division by zero  
        diff = self.max_val - self.min_val  
        epsilon = 1e-7  
        self.max_val[diff < epsilon] = \  
            self.min_val[diff < epsilon] +  
            epsilon  
    def normalize(self, x):  
        return (x - self.min_val) / \  
            (self.max_val - self.min_val)  
  
# Coordinate transformation to polar  
if COORDINATES == "1":  
    u_squared = tensor[0, :, :]**2  
    v_squared = tensor[1, :, :]**2  
    magnitude = torch.sqrt(u_squared + v_squared)  
    )  
    direction = (180 + (180/math.pi) *  
                torch.atan2(-u_squared, -  
                            v_squared)) % 360  
    tensor = torch.stack([magnitude, direction],  
                        axis=0)
```

# Model and hyperparameters - General differences



Interesting points in our implementation:

- Conv2d with stride=2 for downsampling.
- BatchNorm to reduce internal covariate shift. This improves speed and stability of the training.

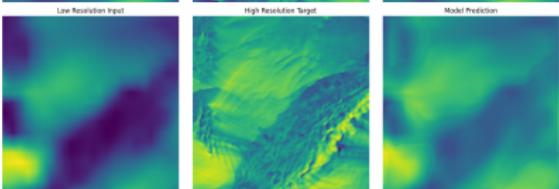
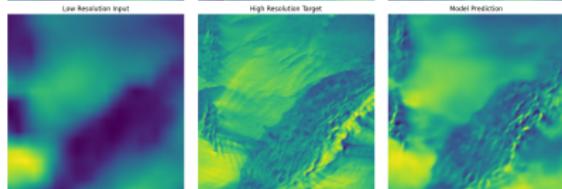
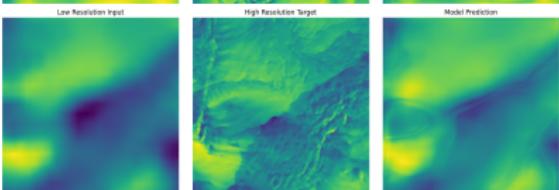
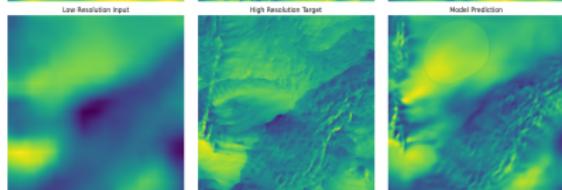
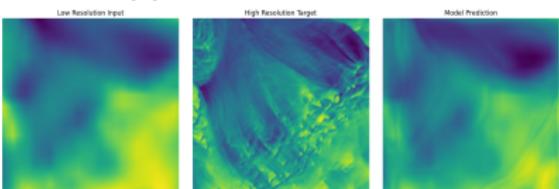
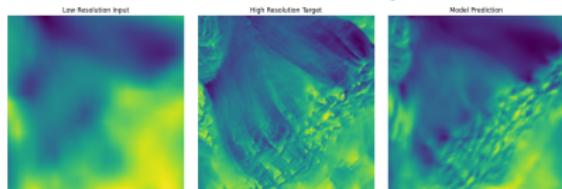
But there's more...

Network Type	Parameters Count	Training Time	Test Loss	Test SSIM
Transposed Conv	13,041,922 (13.0 M)	54:03 16.22s/it	0.19775	0.70910
Bilinear Interp	3,607,586 (3.6 M)	08:04 2.42s/it	0.17166	0.74983

**Table:** Comparison of bilinear interpolation vs transposed convolution.

# Model and hyperparameters - Loss

Even the best model couldn't get decent results with the suggested MSE loss, so we adopted a combined loss approach



L1 + SSIM loss. SSIM: 0.7556

MSE loss. SSIM: 0.7022

## Model and hyperparameters - Loss

```
class L1SSIMLoss(nn.Module):
    def __init__(self, alpha=0.85, ssim_window_size=11, ssim_data_range=1.0, ssim_channel=1):
        super(L1SSIMLoss, self).__init__()
        self.alpha = alpha
        self.l1_loss = nn.L1Loss() # Mean Absolute Error
        self.ssim_loss_fn = SSIMLoss(window_size=ssim_window_size, data_range=ssim_data_range, channel=ssim_channel)

    def forward(self, y_pred, y_true):
        ssim_val_loss = self.ssim_loss_fn(y_pred, y_true)
        l1_val_loss = self.l1_loss(y_pred, y_true)

        combined_loss = self.alpha * ssim_val_loss + (1 - self.alpha) * l1_val_loss
        return combined_loss
```

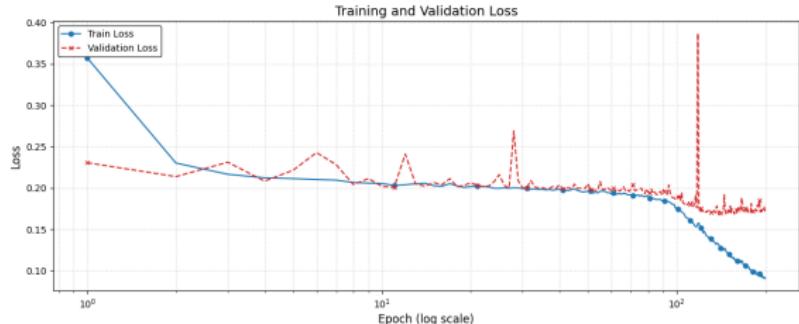
# Model and hyperparameters - Hyperparameters

Channels	Epochs	Loss	Batch Size	LR	LR Scheduler	Weight Decay	Test Loss	Test SSIM
[32, 64, 128]	200	L1 + SSIM	8	1.00E-03	NO	1.00E-05	0.19090	0.71990
[16, 32, 64, 128]	200	L1 + SSIM	8	1.00E-03	NO	1.00E-05	0.17700	0.74170
[16, 32, 64, 128, 256]	300	L1 + SSIM	8	1.00E-03	NO	1.00E-05	<b>0.16768</b>	<u>0.75555</u>
[16, 32, 64, 128, 256]	300	MSE + SSIM	8	1.00E-03	NO	1.00E-05	<u>0.15464</u>	<u>0.75391</u>
[16, 32, 64, 128, 256]	300	MSE	8	1.00E-03	NO	1.00E-05	0.00327	0.70218
[16, 32, 64, 128, 256, 512]	300	L1 + SSIM	8	1.00E-03	NO	1.00E-05	0.19207	0.71774
[16, 32, 64, 128, 256, 512]	300	L1 + SSIM	8	1.00E-03	YES	1.00E-05	0.19207	0.71774

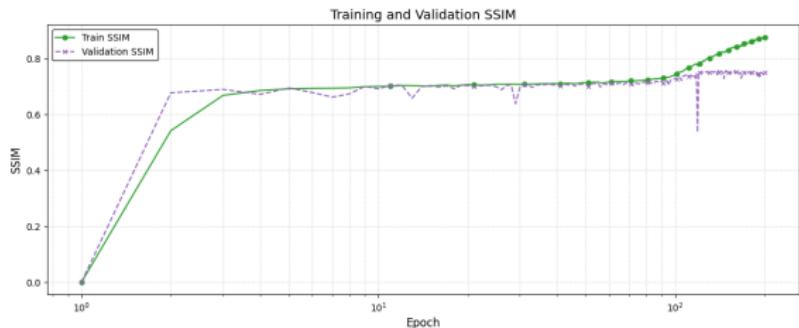
**Table:** Performance comparison of different network configurations. Best and runner-up models are highlighted in bold and underline, respectively.

# Local Training - Loss and SSIM curves

Training curves for the best model on local data (loss and SSIM):

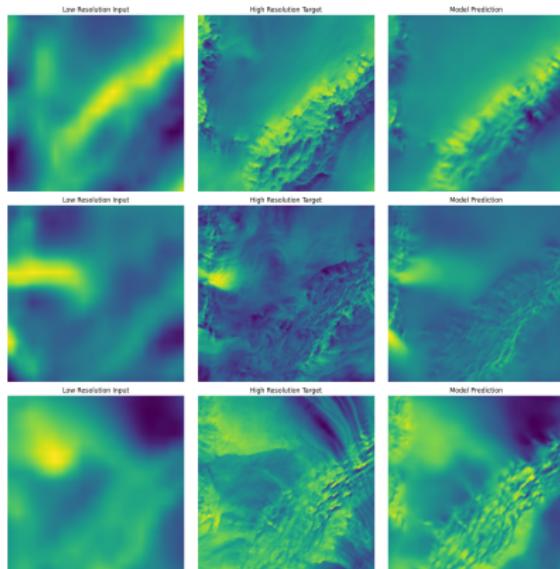


Loss curve

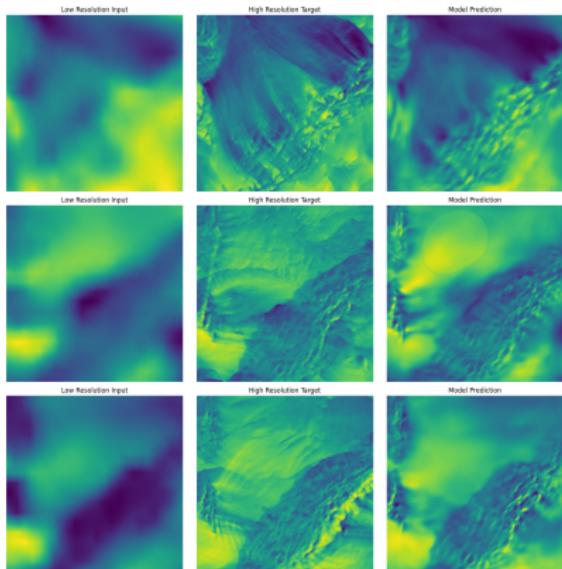


SSIM curve

# Colab Training - Difference with Local data



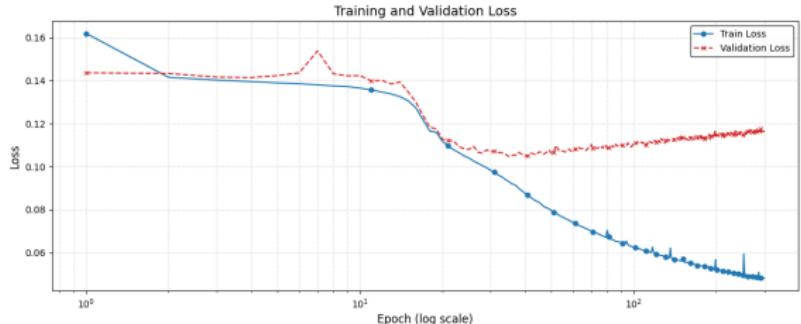
Colab data. SSIM: 0.84483



Local data. SSIM: 0.7556

# Colab Training - Loss and SSIM curves

Training curves for the best model on colab data (loss and SSIM):

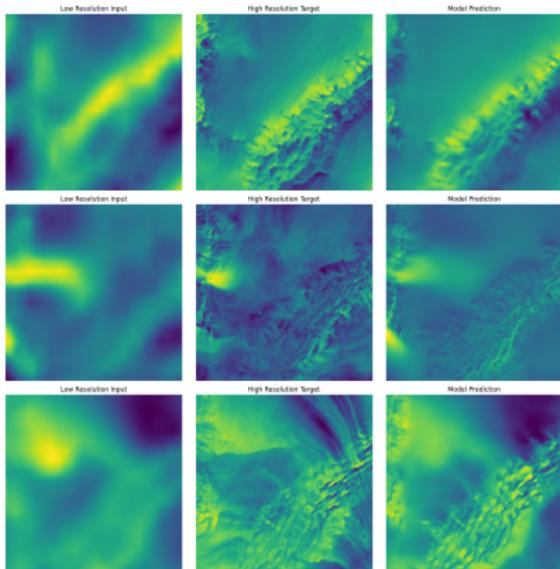


Loss curve

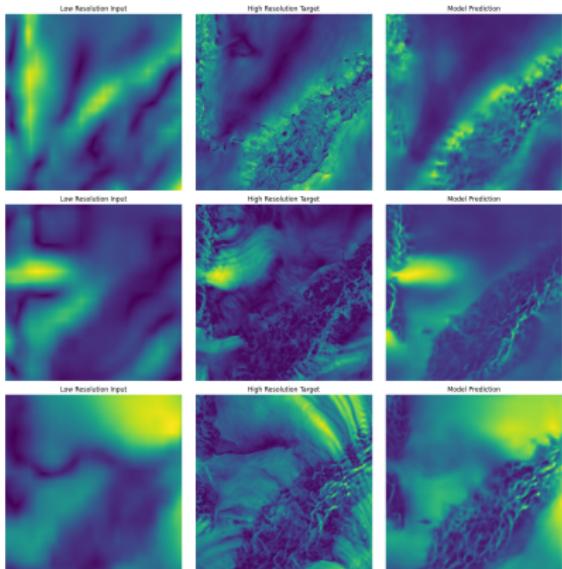


SSIM curve

# Colab Training - Vectorial data vs Directional data



Vector data. SSIM: 0.84483



Direction data. SSIM: 0.48154

# Alternative Datasets and Loss Function Experiments

Dataset	Loss	Observation	Test Loss	SSIM
Local direction	L1 + SSIM	Poor reconstruction performance	0.4912	0.3896
Local direction	MSE	Slight improvement, but low SSIM	0.0480	0.3563
<b>Colab vectors</b>	L1 + SSIM	Best overall performance	<b>0.1052</b>	<b>0.8445</b>
Local vectors	MSE + SSIM	Balanced compromise	0.1546	0.7539
Local vectors	MSE	Suboptimal results	0.3949	0.5229

- **Insight:** Directional data underperformed vs. vector-based representations.
- Combined losses (e.g., L1 + SSIM) improved perceptual quality.

## Societal Impact

More accurate high-resolution forecasts support *agricultural management, climate disaster preparedness and the coastal/tourist community.*

## Next steps

- Check whether the  $L1 + SSIM$  loss still favors accurate reconstruction in complex areas with clear structures, compared to flatter regions where features are harder to predict.
- If not, consider analyzing the alpha weight in the loss function to see if the balance between L1 and SSIM is skewed, causing the network to rely too heavily on one component.

- [1] Fabio Merizzi, Andrea Asperti, and Stefano Colamonaco. "Wind speed super-resolution and validation: from ERA5 to CERRA via diffusion models". In: *Neural Computing and Applications* 36.34 (2024), pp. 21899–21921.
- [2] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. "U-net: Convolutional networks for biomedical image segmentation". In: *Medical image computing and computer-assisted intervention–MICCAI 2015: 18th international conference, Munich, Germany, October 5–9, 2015, proceedings, part III* 18. Springer. 2015, pp. 234–241.
- [3] Robbie A Watt and Laura A Mansfield. "Generative diffusion-based downscaling for climate". In: *arXiv preprint arXiv:2404.17752* (2024).